


⋮



Week 8 Discussion

[Dave DeBarr](#)

[All Sections](#)

4 11

Subscribe"/>

[← Reply](#)

O

[Luis Garduno](#)

May 26, 2021

⋮

Hi Dave,

I was able to run the code for this week's homework on the Azure VM without issues, but because I wanted to play with bigger models, I setup my own physical Ubuntu machine with PyTorch and I'm getting an error when the Python script runs the "trainer.train()" line.

Here's the output and error I'm getting:

```
(dp) luis@ubuntu:~/deep-learning-uw/lesson-8$ time python imdb-train.py roberta-base
2021-05-26 19:32:36.215611: I tensorflow/stream_executor/platform/default/dso_loader.cc:49] Successfully opened dynamic library libc
Some weights of the model checkpoint at roberta-base were not used when initializing RobertaForSequenceClassification: ['roberta.pooler.dense.weight', 'roberta.pooler.dense.bias']
- This IS expected if you are initializing RobertaForSequenceClassification from the checkpoint of a model trained on another task or with a different configuration file.
- This IS NOT expected if you are initializing RobertaForSequenceClassification from the checkpoint of a model that you expect to be initialized from without needing any post-processing.
Some weights of RobertaForSequenceClassification were not initialized from the model checkpoint at roberta-base and are newly initialized from random normal distribution.
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
0%|          | 0/3125 [00:00<?, ?it/s]Traceback (most recent call last):
  File "imdb-train.py", line 103, in <module>
    trainer.train()
  File "/home/luis/anaconda3/envs/dp-uw/lib/python3.8/site-packages/transformers/trainer.py", line 1272, in train
    tr_loss += self.training_step(model, inputs)
  File "/home/luis/anaconda3/envs/dp-uw/lib/python3.8/site-packages/transformers/trainer.py", line 1752, in training_step
    loss.backward()
  File "/home/luis/anaconda3/envs/dp-uw/lib/python3.8/site-packages/torch/tensor.py", line 245, in backward
    torch.autograd.backward(self, gradient, retain_graph, create_graph, inputs=inputs)
  File "/home/luis/anaconda3/envs/dp-uw/lib/python3.8/site-packages/torch/autograd/__init__.py", line 145, in backward
    Variable._execution_engine.run_backward(
RuntimeError: CUDA error: CUBLAS_STATUS_EXECUTION_FAILED when calling `cublasSgemm( handle, opa, opb, m, n, k, &alpha, a, lda, b, ldb, 0, 0, 0, 0, CUBLAS_WORKSPACE_CONFIG=<as-is> )`
real 0m25.954s
user 1m24.932s
sys 0m6.330s
```

Would this be due to something I missed while configuring my machine? I searched online but couldn't figure it out, so any pointers will be appreciated :)

Oh, an these are the versions of CUDA and PyTorch on my machine, in case it helps:

CUDA Version: 11.3

| | |
|-------------|-------------|
| torch | 1.8.1+cu111 |
| torchaudio | 0.8.1 |
| torchvision | 0.9.1+cu111 |

The only thing I see different is that my machine has newer versions of CUDA and PyTorch, and that I don't have the torchtext package installed, but I don't think that would be the reason.

Thanks,

Luis

Edited by [Luis Garduno](#) on May 26 at 10:18pm

[← Reply](#) 



[Luis Garduno](#)

May 27, 2021

Ok, I think my earlier issue was due to having installed both tensorflow and pytorch in the same conda environment.

I created a new environment and only installed pytorch there, but now I'm getting another error... < sigh >

```
(conda-pytorch) luis@ubuntu:~/deep-learning-uw/lesson-8$ python imdb-train.py roberta-base
Traceback (most recent call last):
  File "imdb-train.py", line 11, in
    from transformers import AutoModelForSequenceClassification, AutoTokenizer, EvalPrediction, Trainer, TrainingArguments
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/site-packages/transformers-4.6.1-py3.8.egg/transformers/__init__.py"
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/site-packages/transformers-4.6.1-py3.8.egg/transformers/file_utils.py"
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/site-packages/transformers-4.6.1-py3.8.egg/transformers/__init__.py"
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/importlib/__init__.py", line 127, in import_module
    return _bootstrap._gcd_import(name[level:], package, level)
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/site-packages/transformers-4.6.1-py3.8.egg/transformers/models/__init__.py"
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/site-packages/transformers-4.6.1-py3.8.egg/transformers/models/layoutlm.py"
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/site-packages/transformers-4.6.1-py3.8.egg/transformers/models/layoutlm.py"
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/site-packages/transformers-4.6.1-py3.8.egg/transformers/models/bert.py"
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/site-packages/transformers-4.6.1-py3.8.egg/transformers/tokenization_utils.py"
  File "/home/luis/anaconda3/envs/conda-pytorch/lib/python3.7/site-packages/tokenizers/__init__.py", line 79, in
    from .tokenizers import (
ImportError: /lib/x86_64-linux-gnu/libm.so.6: version `GLIBC_2.29' not found (required by /home/luis/anaconda3/envs/conda-pytorch/
```

I'll have to debug this one later.

[← Reply](#) 



[Dave DeBarr](#)

May 27, 2021

For this first one, it looks like the wrong version of the CUDA toolkit? For example, at the top it says, "Successfully opened dynamic library lib cudart.so.10.1"; but then at the bottom you say you're using CUDA 11.

I use CUDA toolkit 11.0 on my machine, and this is what it says when starting a job, "Successfully opened dynamic library lib cudart.so.11.0" (e.g. <https://www.cross-entropy.net/ML530/imdb-bert-models.txt>).

[← Reply](#)  (1 like)



[Dave DeBarr](#)

May 27, 2021

How are you creating your environment, and which version of python are you using?

For what it's worth: I'm using python version 3.8.5 on my computer ("python --version").

[← Reply](#) 



[Dave DeBarr](#)

May 27, 2021

This person may have run into the same problem you encountered?

<https://discuss.huggingface.co/t/problem-installing-using-conda/5518>

[← Reply](#)  (1 like)



[Luis Garduno](#)

May 27, 2021

Ah, thanks for pointing me to that post! The suggestion there to use the following command fixed my problem:

```
conda install -c conda-forge transformers
```

instead of the recommended command from huggingface.co:

```
conda install -c huggingface transformers
```

To answer your other questions, I followed your setup steps from <https://www.cross-entropy.net/ML530/setup.txt> (minus installing tensorflow), and I created my conda environment using python 3.8, too.

I then installed pytorch with the following command:

```
conda install pytorch torchvision torchtex torchaudio cudatoolkit=11.1 -c pytorch -c nvidia
```

And finally used the command to install transformers mentioned above.

Thanks again,
Luis

← [Reply](#) 👍



[Dave DeBarr](#)

May 27, 2021

Glad it worked!

Just double checking: when you execute "nvcc --version" you see 11.1?

It's funky that it was finding cuda 10.1 earlier; but hopefully all is well now.

← [Reply](#) 👍

[SHOW MORE](#)

1 1



[Roman Battisti](#)

May 31, 2021

Hi Dave,

I'm working on the homework and get the following errors:

```
(py37_pytorch) deeplearning@ML-RefVm-871628:~$ python imdb-train.py.txt rober
```

ta-base

Traceback (most recent call last):

```
File "imdb-train.py.txt", line 11, in <module>
```

```
    from transformers import AutoModelForSequenceClassification, AutoTokenizer
    EvalPrediction, Trainer, TrainingArguments
```

r,

```
File "/home/deeplearning/transformers.py", line 1, in <module>
```

```
    import tensorflow as tf
```

```
ModuleNotFoundError: No module named 'tensorflow'
```

So I pip installed tensorflow. Then I get the following error:

```
(py37_pytorch) deeplearning@ML-RefVm-871628:~$ python imdb-train.py.txt rober
```

ta-base

```
2021-05-31 20:55:44.225669: W tensorflow/stream_executor/platform/default/dso
```

_loader.cc:64]

```
Could not load dynamic library 'libcudart.so.11.0'; dlderror: I
```

libcudart.so.11.0: cannot open shared

```
object file: No such file or directory;
```

LD_LIBRARY_PATH:

```
/opt/intel/compilers_and_libraries_2018.1.163/linux/tbb/lib/
```

```
intel64_lin/gcc4.7:/opt/intel/compilers_and_libraries_2018.1.163/linux/compil
```

```
er/lib/intel64_lin:/opt/intel/compilers_and_libraries_2018.1.163/linux/mkl/li
```

```
b/intel64_lin:/usr/local/cuda/lib64:/usr/local/cuda/extras/CUPTI/lib64/
```

2021-05-31 20:55:44.225709: I tensorflow/stream_executor/cuda/cudart_stub.cc:
above cudart dlerror if you do not have a GPU set up on your machi
Traceback (most recent call last):
File "imdb-train.py.txt", line 11, in <module>
from transformers import AutoModelForSequenceClassification, AutoTokenizer
EvalPrediction, Trainer, TrainingArguments
ImportError: cannot import name 'AutoModelForSequenceClassification' from 'tr
(/home/deeplearning/transformers.py)

29] Ignore

ne.

r,

ansformers'

Any suggestions? Thank you for the help.

← [Reply](#) 👍

[Dave DeBarr](#)

May 31, 2021

I would not have voted for using "pip install tensorflow".

Did you remember to "pip install transformers"?

I'm thinking you were in a directory with transformers.py and so it thought it should try to use the transformers.py module from last week's homework.

Maybe create a separate directory for this week's assignment (imdb)?

Edited by [Dave DeBarr](#) on May 31 at 2:37pm

← [Reply](#) 👍

[Roman Battisti](#)

May 31, 2021

I uninstalled tensorflow and created a new directory for the imdb data. That seems to have worked. Thank you again for the help.

← [Reply](#) 👍 (1 like)